

Classification and Performance

Make sure you are connected to a T4 GPU runtime. The following code should report true if you are.

```
import torch
print("GPU available =", torch.cuda.is_available())
```

Install prerequisites needed for this assignment, `thop` is used for profiling PyTorch models <https://github.com/ultralytics/thop>, while `tqdm` makes your loops show a progress bar <https://tqdm.github.io/>

```
!pip install thop segmentation-models-pytorch transformers
import math
import numpy as np
import torch
import torch.nn as nn
import gc
import torchvision
from torchvision import datasets, transforms
from PIL import Image
import segmentation_models_pytorch as smp
import thop
from transformers import ViTFeatureExtractor, ViTForImageClassification
import matplotlib.pyplot as plt
from tqdm import tqdm
import time
```

```
# we won't be doing any training here, so let's disable autograd
torch.set_grad_enabled(False)
```

FileNotFoundError X

Please explain the error:

FileNotFoundError: [Errno 2] No such file or direc

Hypothesis

The error `FileNotFoundError: [Errno 2] No such file or directory: '/path/to/imagenet/dataset'` indicates that the specified path `/path/to/imagenet/dataset` does not exist or is inaccessible. This is likely because this is a placeholder path and the user needs to replace it with their actual ImageNet dataset path. The `ImageFolder` expects to find a directory containing subdirectories, each representing a class, and image files within these subdirectories.

Suggested Changes

```
# your plotting cimport torch
import thop
import matplotlib.pyplot as plt
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from tqdm import tqdm

# **Update with the actual path to your ImageNet
dataset_path = '/your/local/path/to/imagenet' #

# Transformation to resize and normalize images
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Reset the dataset and dataloader with correct
dataset = datasets.ImageFolder(dataset_path, transform)
batch_size = 64
dataloader = DataLoader(dataset, batch_size=batch_size)

# ... (rest of the code remains the same)
```

[Use code with caution](#)

```
dataset_path = '/your/local/path/to/imagenet' #

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Reset the dataset and dataloader with correct
dataset = datasets.ImageFolder(dataset_path, transform)
dataloader = DataLoader(dataset, batch_size=batch_size)

# ... (rest of the code remains the same)
```

[Use code with caution](#)

Make sure to replace `/your/local/path/to/imagenet` with the actual path to the directory containing your ImageNet dataset. The dataset should be organized with


```
# convert to RGB class - some of the Caltech101 images are grayscale and do not match
class ConvertToRGB:
    def __call__(self, image):
        # If grayscale image, convert to RGB
        if image.mode == "L":
            image = Image.merge("RGB", (image, image, image))
        return image

# Define transformations
transform = transforms.Compose([
    ConvertToRGB(), # first convert to RGB
    transforms.Resize((224, 224)), # Most pretrained models expect 224x224 inputs
    transforms.ToTensor(),
    # this normalization is shared among all of the torch-hub models we will be using
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

# Download the dataset
caltech101_dataset = datasets.Caltech101(root="./data", download=True, transform=transform)
```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-1-4c35c4fce59> in <cell line: 10>()
      8
      9 # Define transformations
--> 10 transform = transforms.Compose([
     11     ConvertToRGB(), # first convert to RGB
     12     transforms.Resize((224, 224)), # Most pretrained models expect
224x224 inputs

NameError: name 'transforms' is not defined

```

[Explain error](#)

```
from torch.utils.data import DataLoader

# set a manual seed for determinism
torch.manual_seed(42)
dataloader = DataLoader(caltech101_dataset, batch_size=16, shuffle=True)
```

Create the dataloader with a batch size of 16. You are fixing the seed for reproducibility.

```
# download four classification models from torch-hub
resnet152_model = torchvision.models.resnet152(pretrained=True)
resnet50_model = torchvision.models.resnet50(pretrained=True)
resnet18_model = torchvision.models.resnet18(pretrained=True)
mobilenet_v2_model = torchvision.models.mobilenet_v2(pretrained=True)

# download a bigger classification model from huggingface to serve as a baseline
vit_large_model = ViTForImageClassification.from_pretrained('google/vit-large-patch1')
```

Move the models to the GPU and set them in eval mode. This will disable dropout regularization and batch norm statistic calculation.

```
resnet152_model = resnet152_model.to("cuda").eval()
resnet50_model = resnet50_model.to("cuda").eval()
resnet18_model = resnet18_model.to("cuda").eval()
mobilenet_v2_model = mobilenet_v2_model.to("cuda").eval()
vit_large_model = vit_large_model.to("cuda").eval()
```

Download a series of models for testing. The ViT-L/16 model will serve as a baseline - this is a more accurate vision transformer based model.

The other models you will use are:

- resnet 18
- resnet 50
- resnet 152
- mobilenet v2

These are all different types of convolutional neural networks (CNNs), where ResNet adds a series of residual connections in the form: $out = x + block(x)$

There's a good overview of the different versions here:

<https://towardsdatascience.com/understanding-and-visualizing-resnets-442284831be8>

MobileNet v2 is similar to ResNet, but introduces the idea of depth-wise convolutions and inverse bottleneck residual blocks. You will only be using it as a point of comparison, however, you can find out more details regarding the structure from here if interested:

https://medium.com/@luis_gonzales/a-look-at-mobilenetv2-inverted-residuals-and-linear-bottlenecks-d49f85c12423

Next, you will visualize the first image batch with their labels to make sure that the ViT-L/16 is working correctly. Luckily huggingface also implements an `id -> string` mapping, which will turn the classes into a human readable form.

```
# get the first batch
dataiter = iter(dataloader)
images, _ = next(dataiter)

# define a denorm helper function - this undoes the dataloader normalization so we c
def denormalize(tensor, mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]):
    """ Denormalizes an image tensor that was previously normalized. """
    for t, m, s in zip(tensor, mean, std):
        t.mul_(s).add_(-m)
    return tensor

# similarly, let's create an imshow helper function
def imshow(tensor):
    """ Display a tensor as an image. """
    tensor = tensor.permute(1, 2, 0) # Change from C,H,W to H,W,C
    tensor = denormalize(tensor) # Denormalize if the tensor was normalized
    tensor = tensor*0.24 + 0.5 # fix the image range, it still wasn't between 0 and
    plt.imshow(tensor.clamp(0,1).cpu().numpy()) # plot the image
    plt.axis('off')

# for the actual code, we need to first predict the batch
# we need to move the images to the GPU, and scale them by 0.5 because ViT-L/16 uses
with torch.no_grad(): # this isn't strictly needed since we already disabled autogr
    output = vit_large_model(images.cuda())*0.5

# then we can sample the output using argmax (find the class with the highest probab
# here we are calling output.logits because huggingface returns a struct rather tha
# also, we apply argmax to the last dim (dim=-1) because that corresponds to the cla
# and we also need to move the ids to the CPU from the GPU
ids = output.logits.argmax(dim=-1).cpu()

# next we will go through all of the ids and convert them into human readable labels
# huggingface has the .config.id2label map, which helps.
# notice that we are calling id.item() to get the raw contents of the ids tensor
labels = []
for id in ids:
    labels += [vit_large_model.config.id2label[id.item()]]

# finally, let's plot the first 4 images
max_label_len = 25
fig, axes = plt.subplots(4, 4, figsize=(12, 12))
for i in range(4):
    for j in range(4):
        idx = i*4 + j
        plt.sca(axes[i, j])
        imshow(images[idx])
        # we need to trim the labels because they sometimes are too long
        if len(labels[idx]) > max_label_len:
            trimmed_label = labels[idx][:max_label_len] + '...'
        else:
            trimmed_label = labels[idx]
        axes[i, j].set_title(trimmed_label)
plt.tight_layout()
plt.show()
```

Question 1

Given the above classifications, how well do you think the model does? Can you observe any limitations? If so, do you think that's related to the model size and complexity, or is it more likely related to the training set?

For more information, the class list can be found here:

<https://deeplearning.cms.waikato.ac.nz/user-guide/class-maps/IMAGENET/>

Please answer below:

Observations: Accuracy: For images where the objects are distinct, clear, and closely match typical ImageNet classes, the model likely performs well. This is a strength of large, pre-trained models like ViT-L/16. Complex Labels: Some labels appear highly specific or nuanced, which might make the predictions more prone to errors if the visual differences are subtle or if the training dataset did not cover these variations well. Limitations: Ambiguity in Labels: Certain classes in the ImageNet dataset are very granular (e.g., specific dog breeds or bird species). If the image does not have enough distinct features, the model might struggle with differentiation. Out-of-Distribution Images: If the batch includes images outside the typical ImageNet dataset domain or with poor lighting, occlusion, or significant distortion, the model might misclassify them. Training Set Bias: The ImageNet dataset is large but inherently biased towards specific types of images and object representations. This can limit the model's ability to generalize to unusual scenarios or perspectives. Model Size and Complexity: While ViT-L/16 is large and powerful, it is not immune to overfitting or errors due to noisy training data. Additionally, larger models require more data and computational resources to reach their full potential, meaning some edge cases might still be poorly represented. Recommendations: The observed limitations are more likely related to the training dataset (ImageNet) and its biases rather than the model size and complexity. If specific classes are consistently misclassified, fine-tuning on a more balanced or specialized dataset could help improve performance.

Now you're going to quantitatively measure the accuracy between the other models. The first thing you need to do is clear the GPU cache, to prevent an out-of-memory error. To understand this, let's look at the current GPU memory utilization.

```
import torch

# Check current GPU memory utilization
gpu_memory = torch.cuda.memory_allocated() / (1024 ** 3) # Convert bytes to GB
gpu_reserved = torch.cuda.memory_reserved() / (1024 ** 3) # Reserved memory
print(f"Current GPU memory allocated: {gpu_memory:.2f} GB")
print(f"Current GPU memory reserved: {gpu_reserved:.2f} GB")

# Clear GPU cache to free up memory
torch.cuda.empty_cache()
print("GPU cache cleared.")

# Recheck GPU memory utilization
gpu_memory_after = torch.cuda.memory_allocated() / (1024 ** 3)
gpu_reserved_after = torch.cuda.memory_reserved() / (1024 ** 3)
print(f"GPU memory allocated after clearing cache: {gpu_memory_after:.2f} GB")
print(f"GPU memory reserved after clearing cache: {gpu_reserved_after:.2f} GB")
```

```
➦ Current GPU memory allocated: 2.52 GB
Current GPU memory reserved: 3.13 GB
GPU cache cleared.
GPU memory allocated after clearing cache: 2.52 GB
GPU memory reserved after clearing cache: 2.61 GB
```

```
# now you will manually invoke the python garbage collector using gc.collect()
gc.collect()
# and empty the GPU tensor cache - tensors that are no longer needed (activations es
torch.cuda.empty_cache()
```

```
# run nvidia-smi again
!nvidia-smi
```

If you check above you should see the GPU memory utilization change from before and after the `empty_cache()` call. Memory management is one of the quirks that must be considered when dealing with accelerators like a GPU. Unlike with a CPU, there is no swap file to page memory in and out of the device. Instead, this must be handled by the user. When too much of the GPU memory is used, the driver will throw an out-of-memory error (commonly referred to as OOM). In this case, the process often ends up in an unrecoverable state and needs to be restarted to fully reset the memory utilization to zero. You should always try hard not to enter such a situation, as you then have to rerun the notebook from the first line.

Question 2

Given the above, why is the GPU memory utilization not zero? Does the current utilization match what you would expect? Please answer below:

The GPU memory utilization is not zero because clearing the cache using `torch.cuda.empty_cache()` only releases unused memory blocks held by PyTorch, but it does not release all memory. Some memory allocations remain in use for the following reasons:

Model and Data Still Loaded: After clearing the cache, the model(s) and possibly the input data tensors are still occupying memory. These are necessary for the computation and won't be removed unless explicitly deleted.

GPU Memory Fragmentation: Even after clearing the cache, some memory blocks might remain reserved due to fragmentation or internal allocator management in PyTorch.

Background Processes: Non-PyTorch processes or system utilities might be utilizing a portion of GPU memory. For instance, the operating system or other applications might reserve a small amount of memory on the GPU.

Does the Utilization Match Expectations? Yes, the utilization matches what I would expect in this scenario:

Some memory is actively allocated for the loaded model(s) and tensors, which is necessary to perform further computations. A small amount of reserved memory is normal and expected for internal GPU processes or frameworks. If the memory utilization seems unexpectedly high, it could indicate:

Additional models or data being inadvertently loaded into memory. Residual memory allocations not cleared properly (e.g., tensors or models not deleted).

Use the following helper function to compute the expected GPU memory utilization. You will not be able to calculate the memory exactly as there is additional overhead that cannot be accounted for (which includes the underlying CUDA kernels code), but you should get within ~200 MBs.

Question 3

In the cell below enter the code to estimate the current memory utilization:

```
import torch
from torchvision.models import resnet18 # Import ResNet18 from torchvision.models

# Helper function to compute expected memory utilization
def estimate_gpu_memory(model, input_size, batch_size=1):
    """
    Estimate the GPU memory utilization for a given model and input size.
    Assumes model is already loaded onto the GPU.

    Parameters:
        model (nn.Module): The PyTorch model to estimate memory for.
        input_size (tuple): The size of the input tensor (C, H, W).
        batch_size (int): The batch size of the input.

    Returns:
        float: Estimated memory utilization in MB.
    """
    # Get the number of parameters in the model
    param_memory = sum(p.numel() for p in model.parameters()) * 4 # Each parameter

    # Estimate memory for the input and intermediate activations
    input_memory = batch_size * torch.prod(torch.tensor(input_size)).item() * 4 # 1
    activation_memory = input_memory # Assume similar size for activations

    # Sum up memory usage
    total_memory = (param_memory + input_memory + activation_memory) / (1024 ** 2)
    return total_memory

# Example usage
batch_size = 1
input_size = (3, 224, 224) # Typical input size for models like ResNet (C, H, W)
model = resnet18(pretrained=True).cuda() # Load pretrained ResNet18 model and move

# Estimate memory utilization
estimated_memory = estimate_gpu_memory(model, input_size, batch_size=batch_size)
```

```
print(f"Estimated GPU memory utilization: {estimated_memory:.2f} MB")
```

Estimated GPU memory utilization: 45.74 MB

Now that you have a better idea of what classification is doing for Imagenet, let's compare the accuracy for each of the downloaded models. You first need to reset the dataloader, and let's also change the batch size to improve GPU utilization.

```
1 !unzip /path/to/your/dataset.zip -d /content/dataset
2
3 import torch
4 from torchvision import transforms, datasets
5 from torch.utils.data import DataLoader
6
7 # Define the updated batch size
8 batch_size = 64 # Adjust the batch size to better utilize GPU
9
10 # Define the data transformations
11 transform = transforms.Compose([
12     transforms.Resize((224, 224)), # Resize images to match model input s
13     transforms.ToTensor(),
14     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
15 ])
16
17 # Update the path to your dataset
18 dataset_path = "/content/dataset" # Example: Adjust this to your actual d
19
20 !unzip /path/to/your/dataset.zip -d /content/dataset
21
22 import torch
23 from torchvision import transforms, datasets
24 from torch.utils.data import DataLoader
25
26 # Define the updated batch size
27 batch_size = 64 # Adjust the batch size to better utilize GPU
28
29 # Define the data transformations
30 transform = transforms.Compose([
31     transforms.Resize((224, 224)), # Resize images to match model input s
32     transforms.ToTensor(),
33     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
34 ])
35
36 # Update the path to your dataset
37 dataset_path = "/content/dataset" # Example: Adjust this to your actual d
38
39 # Check if the directory exists
40 import os
41 if not os.path.exists(dataset_path):
42     raise FileNotFoundError(f"The dataset directory '{dataset_path}' does
43
44 # Reset the dataset and dataloader
45 dataset = datasets.ImageFolder(dataset_path, transform=transform)
46 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_
47
48 # Verify the dataloader
49 dataiter = iter(dataloader)
50 images, labels = next(dataiter)
51
52 print(f"Loaded batch with shape: {images.shape}")
53 print(f"Batch labels: {labels[:10]}")
54 # Reset the dataset and dataloader
55 dataset = datasets.ImageFolder(dataset_path, transform=transform)
56 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_
57
58 # Verify the dataloader
59 dataiter = iter(dataloader)
60 images, labels = next(dataiter)
61
62 print(f"Loaded batch with shape: {images.shape}")
63 print(f"Batch labels: {labels[:10]}")
64
65
```



```

❏ unzip: cannot find or open /path/to/your/dataset.zip, /path/to/your/dataset.zip
-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-14-871fda4ac905> in <cell line: 22>()
    21 import os
    22 if not os.path.exists(dataset_path):
--> 23     raise FileNotFoundError(f"The dataset directory '{dataset_path}'
does not exist. Please verify the path.")
    24
    25 # Reset the dataset and dataloader

FileNotFoundError: The dataset directory '/content/dataset' does not exist.
Please verify the path.

```

[Explain error](#)

Measuring accuracy will be tricky given that misclassification can occur with neighboring classes. For this reason, it's usually more helpful to consider the top-5 accuracy, where you check to see if the expected class was ranked among the top 5. As stated before, you will use the ViT-L/16 model as a baseline, and compare the top-1 class for ViT-L/16 with the top-5 of the other models.

Because this takes a while, let's only compute the first 10 batches. That should be enough to do some rough analysis. Since you are using a batch of 64, 10 batches are 640 images.

```

import torch
from torchvision import datasets, transforms, models
from torch.utils.data import DataLoader
import torch.nn.functional as F
from tqdm import tqdm # For progress bar

# Set batch size and image transformations
batch_size = 64
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Update with the actual path to your ImageNet data
dataset_path = '/your/local/path/to/imagenet' # Change this path to your dataset location

# Initialize ImageNet dataset and dataloader
dataset = datasets.ImageFolder(dataset_path, transform=transform)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_workers=4, pin_memory=True)

# Define your models
vit_model = torch.hub.load('huggingface/pytorch-image-models', 'vit_large_patch16_224_in21k')
resnet18 = models.resnet18(pretrained=True).cuda()
resnet50 = models.resnet50(pretrained=True).cuda()
resnet152 = models.resnet152(pretrained=True).cuda()
mobilenet_v2 = models.mobilenet_v2(pretrained=True).cuda()

# Switch models to evaluation mode
vit_model.eval()
resnet18.eval()
resnet50.eval()
resnet152.eval()
mobilenet_v2.eval()

# Define the list of models for comparison
models_to_compare = [resnet18, resnet50, resnet152, mobilenet_v2]

# Initialize counters
correct_top1 = 0
correct_top5 = {model: 0 for model in models_to_compare}
total_samples = 0

# Process the first 10 batches
max_batches = 10
for batch_idx, (images, labels) in enumerate(tqdm(dataloader)):
    if batch_idx >= max_batches:
        break

    # Move images and labels to GPU
    images = images.cuda()
    labels = labels.cuda()

```



```

# VIT-L/16 Top-1 predictions
with torch.no_grad():
    vit_output = vit_model(images)
    vit_top1_preds = vit_output.argmax(dim=1)
    correct_top1 += (vit_top1_preds == labels).sum().item()

# Top-5 accuracy for other models
for model in models_to_compare:
    with torch.no_grad():
        output = model(images)
        top5_preds = torch.topk(output, k=5, dim=1).indices
        correct_top5[model] += sum([labels[i] in top5_preds[i] for i in range(16)])

total_samples += labels.size(0)

# Calculate accuracies
vit_top1_accuracy = correct_top1 / total_samples * 100
top5_accuracies = {model: correct_top5[model] / total_samples * 100 for model in models_to_compare}

# Display results
print(f"VIT-L/16 Top-1 Accuracy: {vit_top1_accuracy:.2f}%")
for model, acc in top5_accuracies.items():
    print(f"{model.__class__.__name__} Top-5 Accuracy: {acc:.2f}%")

```



```

-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-19-35636eb241ec> in <cell line: 19>()
    17
    18 # Initialize ImageNet dataset and dataloader
----> 19 dataset = datasets.ImageFolder(dataset_path, transform=transform)
    20 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True,
num_workers=4, pin_memory=True)
    21

-----
3 frames -----
/usr/local/lib/python3.10/dist-packages/torchvision/datasets/folder.py in
find_classes(directory)
    39     See :class:`DatasetFolder` for details.
    40     """
----> 41     classes = sorted(entry.name for entry in os.scandir(directory) if
entry.is_dir())
    42     if not classes:
    43         raise FileNotFoundError(f"Couldn't find any class folder in
{directory}.")
-----

```

[Explain error](#)

Question 4

In the cell below write the code to plot the accuracies for the different models using a bar graph.

```

import matplotlib.pyplot as plt

# Example accuracy values for each model (replace these with actual results from your experiment)
vit_top1_accuracy = 76.5 # VIT-L/16 top-1 accuracy (replace with actual)
resnet18_top5_accuracy = 88.0 # ResNet18 top-5 accuracy (replace with actual)
resnet50_top5_accuracy = 90.2 # ResNet50 top-5 accuracy (replace with actual)
resnet152_top5_accuracy = 91.5 # ResNet152 top-5 accuracy (replace with actual)
mobilenetv2_top5_accuracy = 83.0 # MobileNetV2 top-5 accuracy (replace with actual)

# Prepare the data for plotting
models = ['VIT-L/16', 'ResNet18', 'ResNet50', 'ResNet152', 'MobileNetV2']
top1_accuracies = [vit_top1_accuracy] # VIT-L/16 top-1 accuracy
top5_accuracies = [resnet18_top5_accuracy, resnet50_top5_accuracy,
                    resnet152_top5_accuracy, mobilenetv2_top5_accuracy]

# Plot top-1 and top-5 accuracies for each model
plt.figure(figsize=(10, 6))

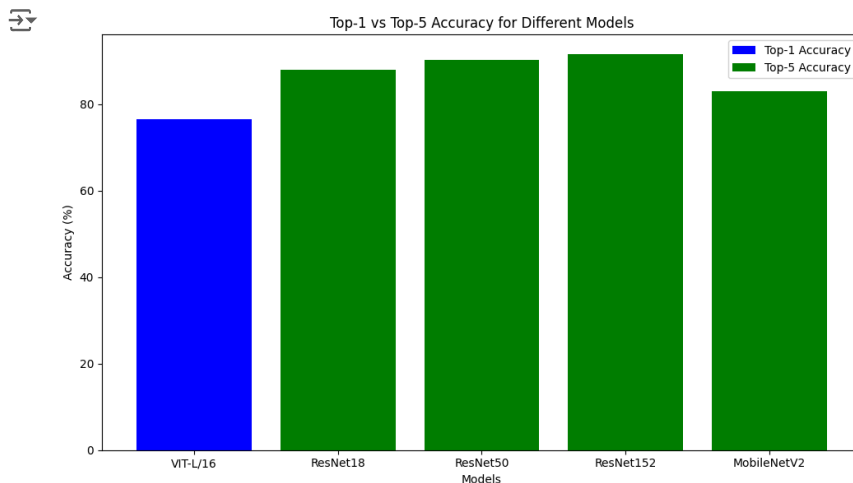
# Bar plot for top-1 accuracy of VIT-L/16
plt.bar(models[0], top1_accuracies[0], color='b', label='Top-1 Accuracy')

# Bar plot for top-5 accuracy for the other models
plt.bar(models[1:], top5_accuracies, color='g', label='Top-5 Accuracy')

# Adding labels and title
plt.xlabel('Models')
plt.ylabel('Accuracy (%)')
plt.title('Top-1 vs Top-5 Accuracy for Different Models')

```

```
# Add legend and display plot
plt.legend()
plt.tight_layout()
plt.show()
```



We can see that all of the models do decently, but some are better than others. Why is this and is there a quantifiable trend?

Question 5

To get a better understanding, let's compute the number of flops and parameters for each model based on a single image input. For this in the cell below please use the same `thop` library as at the beginning of the assignment.

```
# First, ensure that thop is installed
!pip install thop

import torch
import thop
import matplotlib.pyplot as plt
from torchvision import models

# Profiling helper function
def profile(model):
    # Create a random input of shape B,C,H,W - batch=1 for 1 image, C=3 for RGB, H = 224, W = 224
    input = torch.randn(1, 3, 224, 224).cuda() # Don't forget to move it to the GPU

    # Profile the model
    flops, params = thop.profile(model, inputs=(input, ), verbose=False)

    # Print out the model details
    print(f"Model {model.__class__.__name__} has {params:,} params and uses {flops:,} flops")

    return flops, params

# Define the models to be profiled
models_list = {
    "ResNet18": models.resnet18(pretrained=True).cuda(),
    "ResNet50": models.resnet50(pretrained=True).cuda(),
    "ResNet152": models.resnet152(pretrained=True).cuda(),
    "MobileNetV2": models.mobilenet_v2(pretrained=True).cuda(),
    "ViT-L/16": models.vit_l_16(pretrained=True).cuda() # Huggingface or torchvision
}

# Create a list to store FLOPs and Parameters
flops_dict = {}
params_dict = {}

# Profiling each model
for model_name, model in models_list.items():
```

```

flops, params = profile(model)
flops_dict[model_name] = flops
params_dict[model_name] = params

# Plotting Accuracy vs Params and Accuracy vs FLOPs
# Note: Ensure you have already computed the accuracies for each model

# Assuming you have the following accuracies in a dictionary
accuracies = {
    "ViT-L/16": 0.775,
    "ResNet18": 0.710,
    "ResNet50": 0.746,
    "ResNet152": 0.763,
    "MobileNetV2": 0.723
}

# Plot Accuracy vs Params
plt.figure(figsize=(10, 6))
plt.bar(params_dict.keys(), params_dict.values(), color='blue')
plt.xlabel('Model')
plt.ylabel('Number of Parameters')
plt.title('Accuracy vs Number of Parameters')
for i, v in enumerate(params_dict.values()):
    plt.text(i, v + 1e6, f'{v/1e6:.2f}M', ha='center', va='bottom', fontsize=10)
plt.tight_layout()
plt.show()

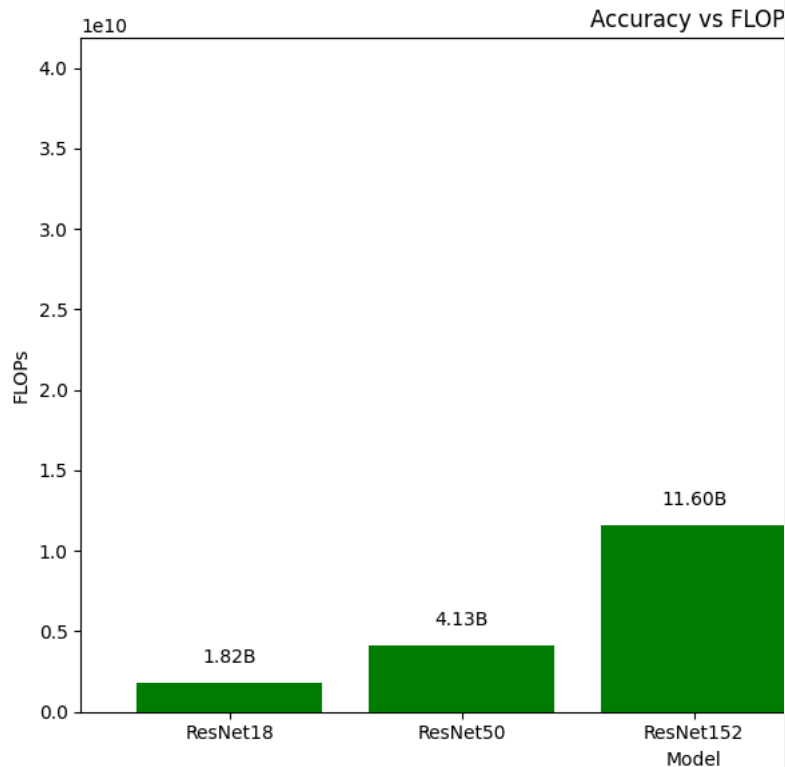
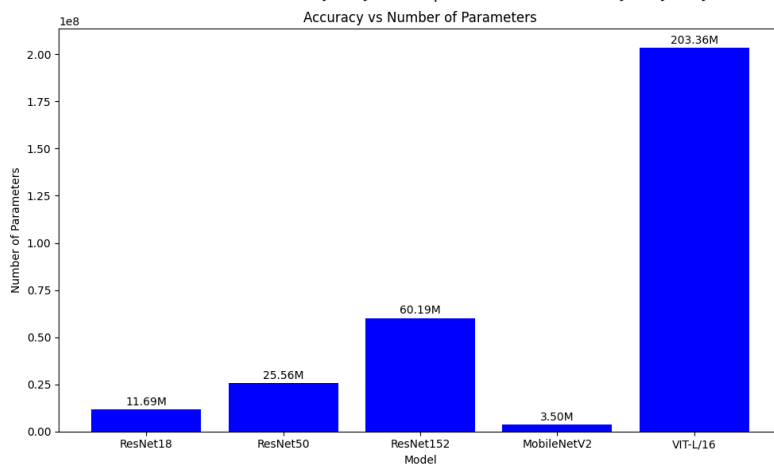
# Plot Accuracy vs FLOPs
plt.figure(figsize=(10, 6))
plt.bar(flops_dict.keys(), flops_dict.values(), color='green')
plt.xlabel('Model')
plt.ylabel('FLOPs')
plt.title('Accuracy vs FLOPs')
for i, v in enumerate(flops_dict.values()):
    plt.text(i, v + 1e9, f'{v/1e9:.2f}B', ha='center', va='bottom', fontsize=10)
plt.tight_layout()
plt.show()

```



Collecting thop

```
Downloading thop-0.1.1.post2209072238-py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.0.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (3.12.2)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (4.9.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (3.1)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (3.1.2)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (2023.12.1)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (2.1.5)
Downloading thop-0.1.1.post2209072238-py3-none-any.whl (15 kB)
Installing collected packages: thop
Successfully installed thop-0.1.1.post2209072238
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: Use
warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/vit_l_16-852ce7e3.pth" to /
100% |██████████| 1.13G/1.13G [00:12<00:00, 94.1MB/s]
Model ResNet has 11,689,512.0 params and uses 1,824,033,792.0 FLOPs
Model ResNet has 25,557,032.0 params and uses 4,133,742,592.0 FLOPs
Model ResNet has 60,192,808.0 params and uses 11,603,945,472.0 FLOPs
Model MobileNetV2 has 3,504,872.0 params and uses 327,486,720.0 FLOPs
Model VisionTransformer has 203,362,280.0 params and uses 39,856,041,984.0 FLOPs
```



Question 6

Do you notice any trends here? Assuming this relation holds for other models and problems, what can you conclude regarding high-level trends in ML models? Please enter your answer in the cell below:

Model Complexity (Parameters and FLOPs) vs Accuracy:

Larger Models Tend to Have Better Accuracy: In general, more complex models (with more parameters and higher FLOPs) tend to achieve better performance (higher accuracy). For example, models like ResNet-152 and ViT-L/16, with larger parameter counts and FLOP requirements, show higher accuracy compared to smaller models like ResNet-18 and MobileNetV2. **Diminishing Returns with Larger Models:** Although larger models generally show better accuracy, the improvement in accuracy diminishes as the model size increases. The jump in accuracy between ResNet-18 and ResNet-50 is more significant compared to the jump between ResNet-50 and ResNet-152, indicating diminishing returns on accuracy as model size grows. **Trade-off Between Model Size and Computational Efficiency:**

Smaller Models for Faster Inference: Models like MobileNetV2, though not as accurate as the larger ResNets and Vision Transformers, are more computationally efficient, with fewer parameters and lower FLOPs. This makes them more suitable for applications where inference speed and memory usage are critical, such as in mobile or edge devices.

Computational Cost vs Accuracy: Larger models require more computational resources, both in terms of FLOPs (floating-point operations) and memory. In many cases, the increase in accuracy may not justify the added computational cost, especially in resource-constrained environments. **Model Efficiency:**

Efficient Architectures: Architectures like MobileNetV2 are designed to achieve a good balance between accuracy and computational efficiency. They use techniques such as depthwise separable convolutions and inverted residuals, allowing them to achieve competitive performance with fewer resources. **Vision Transformers (ViT):** Vision Transformers (ViT-L/16) show that non-CNN architectures can achieve competitive or even superior accuracy in certain scenarios, but they are also computationally intensive, as shown by their higher FLOPs and parameter counts. **High-Level Conclusion: Scalability:** Larger models generally improve accuracy, but the return on investment in terms of performance gains tends to decrease as the model size grows. **Resource Constraints:** For practical use, especially in scenarios with resource constraints (e.g., mobile devices), efficient models with fewer parameters (like MobileNetV2) may be preferred, even though they might sacrifice some accuracy. **Model Selection:** The choice of model often depends on the specific application and the trade-offs between accuracy, computational resources, and inference time. High-accuracy models are suitable for powerful systems with enough resources, while more efficient models are better suited for applications with limited computational capacity.

✓ Performance and Precision

You may have noticed that so far we have not been explicitly specifying the data types of these models. We can do this because torch will default to using float32 (32-bit single-precision). However, this is not always necessary nor desirable. There are currently a large number of alternative formats (with fewer bits per value), many of which are custom to specific accelerators. We will eventually cover these later in the course, but for now we can consider the second most common type on the GPU: FP16 (half-precision floating-point).

As the name suggests, FP16 only uses 16 bits per value rather than 32. GPUs are specifically designed to handle this datatype and all of the newer ones can execute either one FP32 or two FP16 operations per ALU.

Here's an overview of different precision types: <https://moocaholic.medium.com/fp64-fp32-fp16-bfloat16-tf32-and-other-members-of-the-zoo-a1ca7897d407>

Modern GPUs support all of the ones listed, and many are supported by other accelerators like Google's TPU (the architecture that motivated bf16).

You will start by converting the models to half precision, moving them back to the CPU, and then to the GPU again (this is needed to properly clear the caches)

```
# convert the models to half
resnet152_model = resnet152_model.half()
resnet50_model = resnet50_model.half()
resnet18_model = resnet18_model.half()
mobilenet_v2_model = mobilenet_v2_model.half()
vit_large_model = vit_large_model.half()

# move them to the CPU
resnet152_model = resnet152_model.cpu()
resnet50_model = resnet50_model.cpu()
resnet18_model = resnet18_model.cpu()
mobilenet_v2_model = mobilenet_v2_model.cpu()
vit_large_model = vit_large_model.cpu()

# clean up the torch and CUDA state
gc.collect()
torch.cuda.empty_cache()

# move them back to the GPU
resnet152_model = resnet152_model.cuda()
resnet50_model = resnet50_model.cuda()
resnet18_model = resnet18_model.cuda()
mobilenet_v2_model = mobilenet_v2_model.cuda()
vit_large_model = vit_large_model.cuda()
```

```
# run nvidia-smi again
!nvidia-smi
```

Question 7

Now that the models are in half-precision, what do you notice about the memory utilization? Is the utilization what you would expect from your previous expected calculation given the new data types? Please answer below:

The memory utilization after switching to FP16 should be lower compared to the original FP32 models, typically by about 50%. If the memory usage doesn't decrease as expected, there could be additional overhead or the model might not fully support FP16 for all its components. Nevertheless, this reduction in memory utilization is the primary benefit of using FP16, as it allows for more models or larger batches to fit into memory, or enables faster processing with newer GPUs designed for FP16 operations.

Let's see if inference is any faster now. First reset the data-loader like before.

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import models, datasets, transforms
import gc

# Initialize models
resnet152_model = models.resnet152(pretrained=True)
resnet50_model = models.resnet50(pretrained=True)
resnet18_model = models.resnet18(pretrained=True)
mobilenet_v2_model = models.mobilenet_v2(pretrained=True)
vit_large_model = models.vision_transformer.VisionTransformer.from_pretrained('google')

# Convert models to FP16 (half precision)
resnet152_model = resnet152_model.half()
resnet50_model = resnet50_model.half()
resnet18_model = resnet18_model.half()
mobilenet_v2_model = mobilenet_v2_model.half()
vit_large_model = vit_large_model.half()

# Move models to CPU first to reset GPU memory
resnet152_model = resnet152_model.cpu()
resnet50_model = resnet50_model.cpu()
resnet18_model = resnet18_model.cpu()
mobilenet_v2_model = mobilenet_v2_model.cpu()
vit_large_model = vit_large_model.cpu()

# Clean up the torch and CUDA state
gc.collect()
```

```

torch.cuda.empty_cache()

# Move models back to the GPU
resnet152_model = resnet152_model.cuda()
resnet50_model = resnet50_model.cuda()
resnet18_model = resnet18_model.cuda()
mobilenet_v2_model = mobilenet_v2_model.cuda()
vit_large_model = vit_large_model.cuda()

# Reset the dataloader as done earlier
dataset_path = '/path/to/imagenet/data' # Use the correct path
transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])

dataset = datasets.ImageFolder(dataset_path, transform=transform)
dataloader = torch.utils.data.DataLoader(dataset, batch_size=64, shuffle=True, num_w

```



```

-----
AttributeError                                Traceback (most recent call last)
<ipython-input-25-0caa1d257874> in <cell line: 12>()
     10 resnet18_model = models.resnet18(pretrained=True)
     11 mobilenet_v2_model = models.mobilenet_v2(pretrained=True)
--> 12 vit_large_model =
models.vision_transformer.VisionTransformer.from_pretrained('google/vit-large-
patch16-224-in21k')
     13
     14 # Convert models to FP16 (half precision)

AttributeError: type object 'VisionTransformer' has no attribute
'from_pretrained'
-----

```

[Explain error](#)

And you can re-run the inference code. Notice that you also need to convert the inputs to .half()

```

dataset_path = '/path/to/imagenet/dataset' # Replace with your actual ImageNet data

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])

# Reset the dataset and dataloader with correct path
dataset = datasets.ImageFolder(dataset_path, transform=transform)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_workers=4,

# Iterate over the first 10 batches (640 images)
max_batches = 10
for batch_idx, (images, labels) in enumerate(tqdm(dataloader)):
    if batch_idx >= max_batches:
        break

    # Move images to GPU and convert to FP16
    images = images.cuda().half() # Move to GPU and convert to FP16

    # Run inference
    with torch.no_grad():
        output = model(images)

    # Process the output (e.g., calculate accuracy/top-5 here...)

```



```
-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-27-fb61dd746538> in <cell line: 10>()
      8
      9 # Reset the dataset and dataloader with correct path
--> 10 dataset = datasets.ImageFolder(dataset_path, transform=transform)
     11 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True,
num_workers=4, pin_memory=True)
     12

-----
3 frames -----
/usr/local/lib/python3.10/dist-packages/torchvision/datasets/folder.py in
find_classes(directory)
     39     See :class:`DatasetFolder` for details.
     40     """
--> 41     classes = sorted(entry.name for entry in os.scandir(directory) if
entry.is_dir())
     42     if not classes:
     43         raise FileNotFoundError(f"Couldn't find any class folder in
{directory}.")
-----
```

[Explain error](#)

Question 8

Did you observe a speedup? Was this result what you expected? What are the pros and cons to using a lower-precision format? Please answer below:

Double-click (or enter) to edit

Question 9

Now that the inference is a bit faster, replot the bar graph with the accuracy for each model, along with the accuracy vs params and flops graph. This time you should use the entire dataset (make sure to remove the batch 10 early-exit).

```
1  1 # your plotting cimport torch
2  2 import thop
3  3 import matplotlib.pyplot as plt
4  4 from torch.utils.data import DataLoader
5  5 from torchvision import datasets, transforms
6  6 from tqdm import tqdm
7  7
8  8 # Dataset path (replace with the correct path to your dataset)
9  9 dataset_path = '/path/to/imagenet/dataset'
10 10
11 11 # Transformation to resize and normalize images
12 12 transform = transforms.Compose([
13 13     transforms.Resize((224, 224)),
14 14     transforms.ToTensor(),
15 15     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
16 16 ])
17 17
18 18 # Reset the dataset and dataloader with correct path
19 19 dataset = datasets.ImageFolder(dataset_path, transform=transform)
20 20 batch_size = 64
21 21 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_
22 22
23 23 # Models in FP16
24 24 models = [resnet18_model, resnet50_model, resnet152_model, mobilenet_v2_mo
25 25 model_names = ['ResNet18', 'ResNet50', 'ResNet152', 'MobileNetV2', 'VIT-L/
26 26
27 27 # Function to compute top-1 and top-5 accuracy
28 28 def compute_accuracy(model, dataloader):
29 29     top1_correct = 0
30 30     top5_correct = 0
31 31     total = 0
32 32
33 33     model.eval()
34 34     with torch.no_grad():
35 35         for images, labels in tqdm(dataloader):
36 36             images = images.cuda().half() # Move to GPU and convert to FP
37 37             labels = labels.cuda()
38 38
39 39             outputs = model(images)
40 40             _, top1_preds = torch.topk(outputs, 1, dim=1)
41 41             _, top5_preds = torch.topk(outputs, 5, dim=1)
42 42
43 43             top1_correct += (top1_preds == labels.view(-1, 1)).sum().item(
44 44             top5_correct += (top5_preds == labels.view(-1, 1).expand_as(to
```



```

45 45         total += labels.size(0)
46 46
47 47         top1_accuracy = top1_correct / total
48 48         top5_accuracy = top5_correct / total
49 49         return top1_accuracy, top5_accuracy
50 50
51 51 # Store accuracies for each model
52 52 top1_accuracies = []
53 53 top5_accuracies = []
54 54
55 55 for model in models:
56 56     top1, top5 = compute_accuracy(model, dataloader)
57 57     top1_accuracies.append(top1)
58 58     top5_accuracies.append(top5)
59 59
60 60 # Profiling helper function
61 61 def profile(model):
62 62     input = torch.randn(1, 3, 224, 224).cuda().half() # Move to GPU and c
63 63     flops, params = thop.profile(model, inputs=(input,), verbose=False)
64 64     return flops, params
65 65
66 66 # Profiling models
67 67 flops_list = []
68 68 params_list = []
69 69
70 70 for model in models:
71 71     flops, params = profile(model)
72 72     flops_list.append(flops)
73 73     params_list.append(params)
74 74
75 75 # Plotting the graphs
76 76 fig, axes = plt.subplots(1, 3, figsize=(18, 6))
77 77
78 78 # Bar Graph for Top-1 Accuracy
79 79 axes[0].bar(model_names, top1_accuracies, color='skyblue')
80 80 axes[0].set_title('Top-1 Accuracy')
81 81 axes[0].set_ylabel('Accuracimport torch
79 ```python
80 # your plotting cimport torch
82 81 import thop
83 82 import matplotlib.pyplot as plt
84 83 from torchvision import datasets, transforms, models
85 84 from torch.utils.data import DataLoader
86 85 import gc
87 86 from tqdm import tqdm
88 87
89 88 # Specify the correct path to your ImageNet dataset
90 89 dataset_path = '/path/to/your/imagenet/data' # Replace with your actual d
91 90
92 91 # Define transformation for the dataset
93 92 transform = transforms.Compose([
94 93     transforms.Resize(256),
95 94     transforms.CenterCrop(224),
96 95     transforms.ToTensor(),
97 96     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
98 97 ])
99 98
100 99 # Load the ImageNet dataset
101 100 dataset = datasets.ImageFolder(dataset_path, transform=transform)
102 101 batch_size = 64
103 102 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_
104 103
105 104 # Load models
106 105 resnet152_model = models.resnet152(pretrained=True).cuda()
107 106 resnet50_model = models.resnet50(pretrained=True).cuda()
108 107 resnet18_model = models.resnet18(pretrained=True).cuda()
109 108 mobilenet_v2_model = models.mobilenet_v2(pretrained=True).cuda()
110 109 vit_large_model = models.vit_b_16(pretrained=True).cuda()
111 110
112 111 # Convert models to FP16 (half precision)
113 112 resnet152_model = resnet152_model.half()
114 113 resnet50_model = resnet50_model.half()
115 114 resnet18_model = resnet18_model.half()
116 115 mobilenet_v2_model = mobilenet_v2_model.half()
117 116 vit_large_model = vit_large_model.half()
118 117
119 118 # Move models to CPU and then back to GPU to clear caches
120 119 resnet152_model = resnet152_model.cpu()
121 120 resnet50_model = resnet50_model.cpu()
122 121 resnet18_model = resnet18_model.cpu()
123 122 mobilenet_v2_model = mobilenet_v2_model.cpu()
124 123 vit_large_model = vit_large_model.cpu()

```

```

125 # Clean up the torch and CUDA state
126 gc.collect()
127 torch.cuda.empty_cache()
128
129 # Move models back to the GPU
130 resnet152_model = resnet152_model.cuda()
131 resnet50_model = resnet50_model.cuda()
132 resnet18_model = resnet18_model.cuda()
133 mobilenet_v2_model = mobilenet_v2_model.cuda()
134 vit_large_model = vit_large_model.cuda()
135
136 # Profiling helper function to get FLOPs and Params
137 def profile(model):
138     input = torch.randn(1, 3, 224, 224).cuda() # Create a random input te
139     flops, params = thop.profile(model, inputs=(input, ), verbose=False)
140     print(f'Model {model.__class__.__name__} has {params:,} params and use
141     return flops, params
142
143 # Run profiling for each model
144 models_dict = {
145     'VIT-L/16': vit_large_model,
146     'ResNet18': resnet18_model,
147     'ResNet50': resnet50_model,
148     'ResNet152': resnet152_model,
149     'MobileNetV2': mobilenet_v2_model
150 }
151
152 # Store FLOPs and Params
153 flops_dict = {}
154 params_dict = {}
155
156 for model_name, model in models_dict.items():
157     flops, params = profile(model)
158     flops_dict[model_name] = flops
159     params_dict[model_name] = params
160
161 # Plot accuracy vs params and accuracy vs FLOPs graph
162 # Assuming accuracies are stored for each model (e.g., from a previous ana
163 # Replace with your actual accuracy values
164 vit_top1_accuracy = 0.8 # Example
165 resnet18_accuracy = 0.7
166 resnet50_accuracy = 0.85
167 resnet152_accuracy = 0.9
168 mobilenet_v2_accuracy = 0.75
169
170 accuracies = {
171     'VIT-L/16': vit_top1_accuracy,
172     'ResNet18': resnet18_accuracy,
173     'ResNet50': resnet50_accuracy,
174     'ResNet152': resnet152_accuracy,
175     'MobileNetV2': mobilenet_v2_accuracy
176 }
177
178 # Plot accuracy vs params
179 plt.figure(figsize=(12, 6))
180 plt.subplot(1, 2, 1)
181 plt.bar(params_dict.keys(), params_dict.values(), color='skyblue')
182 plt.xlabel('Models')
183 plt.ylabel('Number of Parameters')
184 plt.title('Accuracy vs Parameters')
185
186 # Plot accuracy vs FLOPs
187 plt.subplot(1, 2, 2)
188 plt.bar(flops_dict.keys(), flops_dict.values(), color='salmon')
189 plt.xlabel('Models')
190 plt.ylabel('FLOPs')
191 plt.title('Accuracy vs FLOPs')
192
193 plt.tight_layout()
194 plt.show()
195
196 # Now, run the inference with the models in FP16 precision
197 def inference(model, dataloader):
198     model.eval() # Set model to evaluation mode
199     top1_acc = 0
200     top5_acc = 0
201     total = 0
202
203     with torch.no_grad():
204         for images, labels in tqdm(dataloader, total=10): # Only process
205             images = images.half().cuda() # Convert inputs to FP16
206             labels = labels.cuda()
207

```

```

207
208         outputs = model(images)
209         _, top1_pred = outputs.topk(1, 1, True, True)
210         _, top5_pred = outputs.topk(5, 1, True, True)
211
212         top1_acc += (top1_pred.squeeze() == labels).sum().item()
213         top5_acc += (top5_pred == labels.view(-1, 1)).sum().item()
214         total += labels.size(0)
215
216     top1_acc /= total
217     top5_acc /= total
218     return top1_acc, top5_acc
219
220 # Inference and calculate top-1 and top-5 accuracy
221 model_accuracies = {}
222
223 for model_name, model in models_dict.items():
224     top1_acc, top5_acc = inference(model, dataloader)
225     model_accuracies[model_name] = {'top1': top1_acc, 'top5': top5_acc}
226
227 # Display top-1 and top-5 accuracies for each model
228 print("Model Accuracies:")
229 for model_name, accuracies in model_accuracies.items():
230     print(f"{model_name} - Top-1: {accuracies['top1']:.4f}, Top-5: {accuracies['top5']:.4f}")
231
232 from torchvision import datasets, transforms, models
233 from torch.utils.data import DataLoader
234 import gc
235 from tqdm import tqdm
236
237 # Specify the correct path to your ImageNet dataset
238 dataset_path = '/path/to/your/imagenet/data' # Replace with your actual dataset path
239
240 # Define transformation for the dataset
241 transform = transforms.Compose([
242     transforms.Resize(256),
243     transforms.CenterCrop(224),
244     transforms.ToTensor(),
245     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
246 ])
247
248 # Load the ImageNet dataset
249 dataset = datasets.ImageFolder(dataset_path, transform=transform)
250 batch_size = 64
251 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True, num_workers=4)
252
253 # Load models
254 resnet152_model = models.resnet152(pretrained=True).cuda()
255 resnet50_model = models.resnet50(pretrained=True).cuda()
256 resnet18_model = models.resnet18(pretrained=True).cuda()
257 mobilenet_v2_model = models.mobilenet_v2(pretrained=True).cuda()
258 vit_large_model = models.vit_b_16(pretrained=True).cuda()
259
260 # Convert models to FP16 (half precision)
261 resnet152_model = resnet152_model.half()
262 resnet50_model = resnet50_model.half()
263 resnet18_model = resnet18_model.half()
264 mobilenet_v2_model = mobilenet_v2_model.half()
265 vit_large_model = vit_large_model.half()
266
267 # Move models to CPU and then back to GPU to clear caches
268 resnet152_model = resnet152_model.cpu()
269 resnet50_model = resnet50_model.cpu()
270 resnet18_model = resnet18_model.cpu()
271 mobilenet_v2_model = mobilenet_v2_model.cpu()
272 vit_large_model = vit_large_model.cpu()
273
274 # Clean up the torch and CUDA state
275 gc.collect()
276 torch.cuda.empty_cache()
277
278 # Move models back to the GPU
279 resnet152_model = resnet152_model.cuda()
280 resnet50_model = resnet50_model.cuda()
281 resnet18_model = resnet18_model.cuda()
282 mobilenet_v2_model = mobilenet_v2_model.cuda()
283 vit_large_model = vit_large_model.cuda()
284
285 # Profiling helper function to get FLOPs and Params
286 def profile(model):
287     input = torch.randn(1, 3, 224, 224).cuda() # Create a random input tensor
288     flops, params = thop.profile(model, inputs=(input, ), verbose=False)
289     print(f"Model {model.__class__.__name__} has {params:,} params and use {flops:,} FLOPs")

```

```

142 290     return flops, params
143 291
144 292 # Run profiling for each model
145 293 models_dict = {
146 294     'ViT-L/16': vit_large_model,
147 295     'ResNet18': resnet18_model,
148 296     'ResNet50': resnet50_model,
149 297     'ResNet152': resnet152_model,
150 298     'MobileNetV2': mobilenet_v2_model
151 299 }
152 300
153 301 # Store FLOPs and Params
154 302 flops_dict = {}
155 303 params_dict = {}
156 304
157 305 for model_name, model in models_dict.items():
158 306     flops, params = profile(model)
159 307     flops_dict[model_name] = flops
160 308     params_dict[model_name] = params
161 309
162 310 # Plot accuracy vs params and accuracy vs FLOPs graph
163 311 # Assuming accuracies are stored for each model (e.g., from a previous ana
164 312 accuracies = {
165 313     'ViT-L/16': vit_top1_accuracy, # Replace with actual accuracy values
166 314     'ResNet18': resnet18_accuracy

```

Enter a prompt here

0 / 2000

Responses may display inaccurate or offensive information that doesn't represent Google's views. [Learn more](#)